High-Performance Computing

OpenACC - Convert color image to gray image

Anand Kamble
Department of Scientific Computing
Florida State University

1 Introduction

In this assignment, we are using OpenACC for parallelization. We are using the `pgc++` compiler for compiling this project.

2 Implementation

2.1 Copying Data from host to device

The `#pragma acc data` directive specifies that the `dataBuf` array should be copied to the device before the parallel region, and copied back to the host after the parallel region. Here we are also specifying that we want to copy the whole array by adding `[0:width*height*3]` in the directive.

Listing 1: Grouping Kernel with Shared Memory

```
#pragma acc data copy(dataBuf[0:width*height*3])
{
...
}
```

2.2 Parallelizing the loops

The `#pragma acc parallel loop collapse(2)` directive tells the compiler to parallelize the nested loops across the available device cores.

Listing 2: Grouping Kernel with Shared Memory

```
#pragma acc parallel loop collapse(2)
{
    for(...){
        for(...){
            ...
        }
    }
}
```

2.3 Timing

In this code, we are using `std::chrono` library to measure the time taken by the RGB to grayscale conversion function.

Listing 3: Grouping Kernel with Shared Memory

```
auto start = std::chrono::high_resolution_clock::now();
// RGB to grayscale conversion
auto stop = std::chrono::high_resolution_clock::now();
auto duration = std::chrono::duration_cast<std::chrono::microseconds>(stop -
    start);
printf("Time taken by function: %d ms\n", duration.count() / 1000);
```

3 Results

We can see very good increase in the performance of the program compared to the CPU with minimal effort. Although, CUDA is faster in this case, the time required to program the CUDA is much higher than OpenACC.

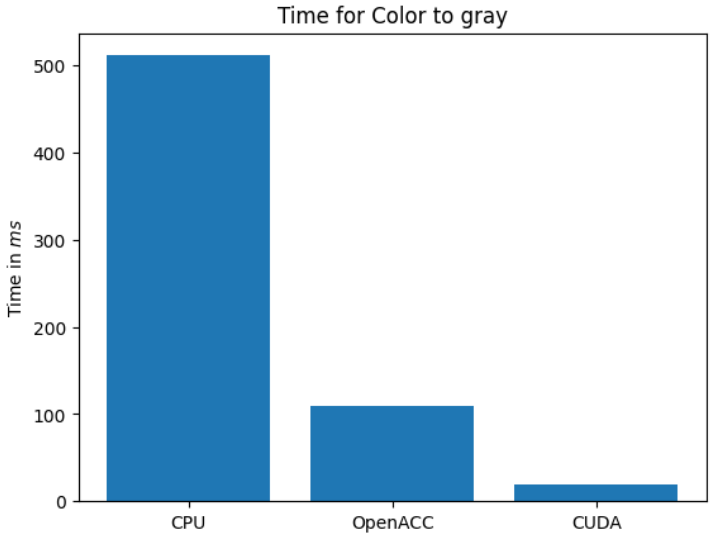

I also did profile the program using nsys, and here are the results. We can see that most of the time spent is for `cuMemAllocHost_v2`.

```
 ** OS Runtime Summary (osrt_sum):

Time (%)  Total Time (ns)  Num Calls    Avg (ns)      Med (ns)     Min (ns)    Max (ns)    StdDev (ns)           Name
--------  ---------------  ---------  -------------  -------------  ---------  -----------  -------------  ----------
    53.9       48,175,787         11  4,379,617.0    913,093.0        1,065  19,626,624    6,817,824.5  poll
    40.4       36,158,173        474     76,283.1      6,121.0          242   9,646,217      620,925.7  ioctl
     1.6        1,393,615         33     42,230.8      2,936.0        1,044   1,194,560      207,244.2  fopen
     1.2        1,069,237         27     39,601.4        860.0          557   1,034,533      198,845.0  fclose
     0.9          841,527         27     31,167.7      4,519.0        3,442     483,519       91,591.5  mmap64
     0.7          599,349          9     66,594.3     20,154.0        8,556     241,538       82,071.0  sem_timedwait
     0.6          536,913          5    107,382.6     73,275.0       49,702     189,197       65,204.6  pthread_create
     0.2          191,612          1    191,612.0    191,612.0      191,612     191,612            0.0  pthread_cond_wait
     0.2          167,568         45      3,723.7      2,891.0        1,533      15,015        2,372.1  open64
     0.1           80,990         15      5,399.3      2,268.0          962      28,954        7,193.2  mmap
     0.0           32,507         49        663.4         36.0           29      30,655        4,373.8  fgets
     0.0           29,230          8      3,653.8      3,317.0          661       6,380        1,948.8  open
     0.0           20,372          4      5,093.0      4,982.5          551       9,856        4,049.1  fread
     0.0           17,903         59        303.4        274.0           93         878          156.7  fcntl
     0.0           14,950         11      1,359.1      1,464.0          421       2,216          596.3  write
     0.0           14,735          6      2,455.8      2,294.0        1,203       4,264        1,181.6  munmap
     0.0           11,774          2      5,887.0      5,887.0        4,860       6,914        1,452.4  socket
     0.0           10,918         15        727.9        525.0          189       1,902          579.9  read
     0.0            9,423          3      3,141.0      3,806.0        1,100       4,517        1,803.0  pipe2
     0.0            9,024          1      9,024.0      9,024.0        9,024       9,024            0.0  fflush
     0.0            7,415          4      1,853.8      1,537.5          126       4,214        1,804.5  fwrite
     0.0            7,316          1      7,316.0      7,316.0        7,316       7,316            0.0  connect
     0.0            3,622          3      1,207.3        303.0          273       3,046        1,592.4  pthread_cond_broadcast
     0.0            1,421          7        203.0        199.0          130         247           37.9  dup
     0.0              972          1        972.0        972.0          972         972            0.0  bind
     0.0              635          1        635.0        635.0          635         635            0.0  listen

Processing [report1.sqlite] with
    [/opt/nvidia/nsight-systems/2023.4.4/host-linux-x64/reports/cuda_api_sum.py]...

** CUDA API Summary (cuda_api_sum):

Time (%)  Total Time (ns)  Num Calls   Avg (ns)     Med (ns)    Min (ns)   Max (ns)   StdDev (ns)          Name
--------  ---------------  ---------  -----------  -----------  ---------  ---------  -----------  ----------
    69.1          518,745          1  518,745.0    518,745.0    518,745    518,745          0.0  cuMemAllocHost_v2
    11.8           88,330          2   44,165.0     44,165.0      6,231     82,099     53,646.8  cuMemAlloc_v2
     6.9           51,872          1   51,872.0     51,872.0     51,872     51,872          0.0  cuModuleLoadDataEx
     4.6           34,538          1   34,538.0     34,538.0     34,538     34,538          0.0  cuMemcpyDtoHAsync_v2
     3.3           25,086          1   25,086.0     25,086.0     25,086     25,086          0.0  cuMemcpyHtoDAsync_v2
     2.6           19,269          1   19,269.0     19,269.0     19,269     19,269          0.0  cuLaunchKernel
     1.6           12,058          3    4,019.3      3,158.0        512      8,388      4,008.0  cuStreamSynchronize
     0.2            1,333          3      444.3        303.0        132        898        402.1  cuCtxSetCurrent
```