

Results

The calculated ratios for $t = 139$ are as follows:

$$\frac{\alpha_{1,139}}{\alpha_{2,139}} = 13.079570428492065$$
$$\frac{\beta_{1,139}}{\beta_{2,139}} = 3.1584763971655754$$

Problem 2: Parameter Estimation using Baum-Welch Algorithm

In this problem, we are provided with a sequence of 20,000 dice rolls (dataset `hmm_pb2.csv`) and tasked with estimating the HMM parameters using the Baum-Welch algorithm.

Initialization

We initialized the initial state probabilities π , transition probabilities a , and emission probabilities b with random values, ensuring they sum to 1.

Baum-Welch Algorithm Implementation

The Baum-Welch algorithm iteratively updated π , a , and b through expectation and maximization steps, converging to estimates for the model parameters.

Results

After 500 iterations, the estimated parameters are:

- **Initial State Probabilities (π):**

$$\pi = [1.000000000e+000 \ 9.61890076e-207]$$

- **State Transition Probabilities (a):**

$$a = \begin{bmatrix} 0.8711917 & 0.1288083 \\ 0.21299416 & 0.78700584 \end{bmatrix}$$

- **Observation Probabilities (b):**

$$b = \begin{bmatrix} 0.21237817 & 0.2137017 & 0.20424693 & 0.18968089 & 0.0900276 & 0.08996471 \\ 0.10482205 & 0.09082342 & 0.11362572 & 0.20883941 & 0.18721835 & 0.29467105 \end{bmatrix}$$

References

- [1] Risto Hinno. *Baum-Welch Algorithm*. Accessed: 2024-10-31. 2023. URL: <https://ristohinno.medium.com/baum-welch-algorithm-4d4514cf9dbe>.
- [2] Abhisek Jana. *Derivation and Implementation of Baum Welch Algorithm for Hidden Markov Model*. Accessed: 2024-10-31. 2019. URL: <https://adeveloperdiary.com/data-science/machine-learning/derivation-and-implementation-of-baum-welch-algorithm-for-hidden-markov-model/>.
- [3] Pierian Training. *Viterbi Algorithm Implementation in Python: A Practical Guide*. Accessed: 2024-10-31. 2023. URL: <https://pieriantraining.com/viterbi-algorithm-implementation-in-python-a-practical-guide/>.

Appendix: Code

The Python code used for the implementation of the Viterbi, forward, backward, and Baum-Welch algorithms is attached below:

```
###
from typing import Any
import numpy as np
import pandas as pd

# Problem 1
# =====

# Read the dataset hmm_pb1.csv
data_pb1: pd.DataFrame = pd.read_csv('hmm_pb1.csv', header=None)
x_pb1: np.ndarray = data_pb1.values.flatten() # Convert to 1D numpy array

# Define the HMM parameters as per the Dishonest casino model
# States: 0='Fair', 1='Loaded'
# Observations: dice rolls (1-6)

# Transition probabilities
a: np.ndarray = np.array([[0.95, 0.05],
                           [0.10, 0.90]])

# Emission probabilities
b: np.ndarray = np.zeros((2, 6))
# Fair die probabilities
b[0, :] = 1/6
# Loaded die probabilities
b[1, :] = [1/10, 1/10, 1/10, 1/10, 1/10, 1/2]

# Initial state probabilities
pi = np.array([0.5, 0.5])

# Convert probabilities to log probabilities to avoid underflow
log_a: np.ndarray = np.log(a)
log_b: np.ndarray = np.log(b)
log_pi: np.ndarray = np.log(pi)

# Implement the Viterbi algorithm using log probabilities
def viterbi_log(x, log_pi, log_a, log_b) -> np.ndarray[Any,
    ↪ np.dtype[np.unsignedinteger[Any]]]:
    N: int = len(log_pi) # Number of states
    T: int = len(x)      # Length of the observation sequence
    delta: np.ndarray[Any, np.dtype[np.float64]] = np.zeros((T, N))
    psi: np.ndarray[Any, np.dtype[Any]] = np.zeros((T, N), dtype=int)

    # Initialization
    delta[0, :] = log_pi + log_b[:, x[0]-1]
    psi[0, :] = 0

    # Recursion
    for t in range(1, T):
        for j in range(N):
            probabilities = delta[t-1, :] + log_a[:, j]
            psi[t, j] = np.argmax(probabilities)
            delta[t, j] = np.max(probabilities) + log_b[j, x[t]-1]
```

```

# Termination
path = np.zeros(T, dtype=int)
path[T-1] = np.argmax(delta[T-1, :])

# Path backtracking
for t in range(T-2, -1, -1):
    path[t] = psi[t+1, path[t+1]]

# Adjust path to match state labels (1='Fair', 2='Loaded')
path += 1
return path

# Run the Viterbi algorithm
y_viterbi = viterbi_log(x_pb1, log_pi, log_a, log_b)

# Output the obtained sequence y for verification
print("Most likely sequence y of states (1='Fair', 2='Loaded'):")
print(y_viterbi)

# Implement the forward algorithm with scaling
def forward_algorithm_scaled(x, pi, a, b) -> tuple[np.ndarray[Any,
↳ np.dtype[np.float64]], np.ndarray[Any, np.dtype[np.float64]]]:
    N: int = len(pi)
    T: int = len(x)
    alpha: np.ndarray[Any, np.dtype[np.float64]] = np.zeros((T, N))
    c: np.ndarray[Any, np.dtype[np.float64]] = np.zeros(T)

    # Initialization
    alpha[0, :] = pi * b[:, x[0]-1]
    c[0] = np.sum(alpha[0, :])
    alpha[0, :] /= c[0]

    # Induction
    for t in range(1, T):
        for j in range(N):
            alpha[t, j] = np.sum(alpha[t-1, :] * a[:, j]) * b[j, x[t]-1]
        c[t] = np.sum(alpha[t, :])
        alpha[t, :] /= c[t]
    return alpha, c

# Implement the backward algorithm with scaling
def backward_algorithm_scaled(x, a, b, c) -> np.ndarray[Any, np.dtype[np.float64]]:
    N = a.shape[0]
    T: int = len(x)
    beta: np.ndarray[Any, np.dtype[np.float64]] = np.zeros((T, N))

    # Initialization
    beta[T-1, :] = 1 / c[T-1]

    # Induction
    for t in range(T-2, -1, -1):
        for i in range(N):
            beta[t, i] = np.sum(a[i, :] * b[:, x[t+1]-1] * beta[t+1, :])
        beta[t, :] /= c[t]
    return beta

# Run the forward and backward algorithms
alpha, c = forward_algorithm_scaled(x_pb1, pi, a, b)

```

```

beta = backward_algorithm_scaled(x_pb1, a, b, c)

# Report alpha_1_139/alpha_2_139 and beta_1_139/beta_2_139
t_index = 139 # t=139
alpha_ratio = alpha[t_index, 0] / alpha[t_index, 1]
beta_ratio = beta[t_index, 0] / beta[t_index, 1]

print("\nalpha_1_139 / alpha_2_139 =", alpha_ratio)
print("beta_1_139 / beta_2_139 =", beta_ratio)

# Problem 2
# =====

# Read the dataset hmm_pb2.csv
data_pb2 = pd.read_csv('hmm_pb2.csv', header=None)
x_pb2 = data_pb2.values.flatten() # Convert to 1D numpy array

# Initialize the HMM parameters randomly
def initialize_random_parameters(N, M):
    pi = np.random.rand(N)
    pi /= np.sum(pi)
    a = np.random.rand(N, N)
    a /= np.sum(a, axis=1, keepdims=True)
    b = np.random.rand(N, M)
    b /= np.sum(b, axis=1, keepdims=True)
    return pi, a, b

# Implement the Baum-Welch algorithm
def baum_welch(x, N, M, max_iter=100):
    T = len(x)
    pi, a, b = initialize_random_parameters(N, M)
    for iteration in range(max_iter):
        # E-step
        alpha, c = forward_algorithm_scaled(x, pi, a, b)
        beta = backward_algorithm_scaled(x, a, b, c)

        log_likelihood = -np.sum(np.log(c))
        print(f"Iteration {iteration+1}, Log-Likelihood: {log_likelihood}")

        gamma = (alpha * beta) / np.sum(alpha * beta, axis=1, keepdims=True)
        xi = np.zeros((T-1, N, N))
        for t in range(T-1):
            denom = np.dot(np.dot(alpha[t, :], a) * b[:, x[t+1]-1], beta[t+1, :])
            for i in range(N):
                numerator = alpha[t, i] * a[i, :] * b[:, x[t+1]-1] * beta[t+1, :]
                xi[t, i, :] = numerator / denom

        # M-step
        pi = gamma[0, :]
        a = np.sum(xi, axis=0) / np.sum(gamma[:-1, :], axis=0)[:, None]
        b_num = np.zeros((N, M))
        for k in range(M):
            indices = np.where(x == k+1)[0]
            if len(indices) > 0:
                b_num[:, k] = np.sum(gamma[indices, :], axis=0)
        b = b_num / np.sum(gamma, axis=0)[:, None]
    return pi, a, b

```

```
# Run the Baum-Welch algorithm on the observed data
N = 2 # Number of states
M = 6 # Number of possible observations (dice rolls 1-6)
pi_estimated, a_estimated, b_estimated = baum_welch(x_pb2, N, M, max_iter=500)

# Report the estimated HMM parameters
print("\nEstimated initial state probabilities (pi):")
print(pi_estimated)

print("\nEstimated state transition probabilities (a):")
print(a_estimated)

print("\nEstimated observation probabilities (b):")
print(b_estimated)
```