

# Introduction to pyMC

## Probabilistic Programming in Python

Sachin Shanbhag

Department of Scientific Computing  
Florida State University,  
Tallahassee, FL 32306.



# Contents

- ▶ Resources
- ▶ Probabilistic Programming
- ▶ pyMC Basics
- ▶ Examples (follow accompanying notebook)
  - ▶ Bayesian Inference to Estimate Cheating Levels
  - ▶ Parameter Estimation and Uncertainty Quantification
  - ▶ Sampling Custom Distributions

pyMC is an active project. Please use a “newer” version (v5)

You can either install pyMC locally or use Google Colab

# Resources

- ▶ Probabilistic Programming in Python with pyMC  
[YouTube video](#) by Chris Fonnesbeck ~ 90 minutes
- ▶ [Introductory Overview](#)  
Based on PeerJ publication, adapted for pyMC v5
- ▶ [pyMC Quickstart Tutorial](#)  
If you want to dig into code directly
- ▶ Visualize Different MCMC Samplers  
[Demo](#)

# Probabilistic Programming

It is a tool for statistical inference.

It is not about writing software that behaves probabilistically.

It is a programming framework in which probabilistic models are specified and inference is performed automatically.

pyMC is one of several probabilistic programming tools

Others include:

Stan	low level, R flavor, market leader
BUGS	standalone, GUI, classic
pyro	Meta, big data, python
Tensorflow probability	Google, python
Edward	python, ML + MCMC
Turing.jl	Julia, ML + MCMC

# Bayesian Inference

Main goal of probabilistic programming is Bayesian inference

$$\pi(\theta|\text{data}) = \frac{\pi(\text{data}|\theta) \times \pi(\theta)}{\pi(\text{data})} \quad (1)$$

If the prior and likelihood are conjugate the posterior distribution  $\pi(\theta|\text{data})$  can be computed analytically

For complex problems, we use MCMC to sample posterior

In probabilistic programming we specify

- ▶ the prior and likelihood as probability distributions
- ▶ the MCMC method to sample the posterior

Framework provides sampling, analysis, and visualization tools.

# pyMC

- ▶ Bayesian inference using MCMC, variational inference etc.
- ▶ [large suite](#) of statistical distributions
- ▶ easy to specify custom distributions that may not be available by default
- ▶ large suite of MCMC algorithms
  - ▶ Metropolis
  - ▶ Gibbs
  - ▶ Hamiltonian Monte Carlo
  - ▶ No U-Turn Sampler
  - ▶ Slice
- ▶ uses [ArviZ](#) for analysis and visualization of the posterior distribution

# Sampling Algorithms

Metropolis MCMC is quite good for small dimensional problems.

But since it is essentially a random walk, it can take a long time to sample large spaces or complex models

I will outline the intuition behind other three popular methods, which shall not otherwise cover in class


- ▶ Slice
- ▶ Hamiltonian Monte Carlo
- ▶ No U-Turn Sampler

Visualization of Samplers

# Slice Sampling

Wikipedia has a helpful [entry](#)

Sophisticated extension of accept/reject



select next sample by uniformly sampling the domain corresponding to a “horizontal slice”



# HMC/NUTS

Consider the energy landscape  $U(\boldsymbol{x})$  corresponding to the target probability distribution  $\pi(\boldsymbol{x})$

$$\pi(\boldsymbol{x}) \sim \exp(-U(\boldsymbol{x})) \quad (2)$$

Hamiltonian Monte Carlo (HMC) essentially simulates the motion of a frictionless “puck” or particle in this landscape

At each step:

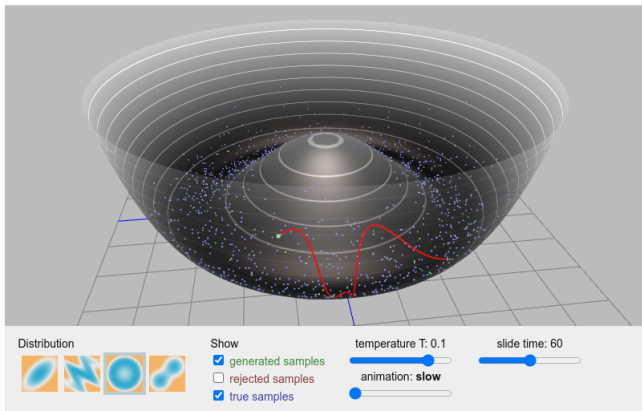
- ▶ sample velocity from normal distribution (kick)
- ▶ find where puck ends up after a certain time interval
- ▶ this location is the new sample

Solving the “equations of motion” using the leap-frog method requires computation of derivatives

pyMC builds computational graphs to compute derivatives

# HMC

Samples are less correlated than Metropolis MCMC



## Animation

Acceptance rates are generally high (of order 0.8)

# NUTS

NUTS (No U-Turn Sampler) is the most common sampling method for continuous variables

It is the default algorithm in pyMC

It is an auto-tuning version of HMC that avoids U-Turns.

What are U-Turns, and why do we want to avoid them?

If you begin climbing a hill from a valley due to a kick, there is a tendency to slip back to the valley during the next turn.

NUTS attempts to avoid this.

# Divergences

Divergences occur when the simulated HMC/NUTS trajectory departs from the true trajectory as measured by total energy

This often occurs when the target distribution has high curvature

Leap-frog takes small steps to simulate particle trajectory.

Small step sizes are inefficient. Large step sizes cause divergences.

pyMC provides warnings about divergences. If there are too many divergences *relative* to the total number of draws you should

- ▶ increase target acceptance rate (longer trajectories)
- ▶ reparametrize the model

# Tutorial

It is best to demonstrate pyMC hands-on

Please see the accompanying Jupyter notebook